

Cloud Suitability Analyzer (CSA) User Manual

Version 2.1.0

Date	Action	Author
Aug 13, 2020	Final edit	S. Woods
Jul 30, 2020	Prep for open source release	S. Woods
Dec 12, 2018	Converted from readme	S. Woods
Dec 18, 2018	Added scoring/graphics	S. Woods
Dec 20, 2018	Added bucketing and profile	S. Woods
Jan 03, 2018	Added EULA	S. Woods

License

Cloud Suitability Analyzer is licensed under `SPDX-License-Identifier: BSD-2`

Purpose

`csa` is built to automatically scan for potential cloud remediation issues (cloud native) and cloud accommodation issues (containerization) embedded in legacy applications. Currently, rules target Java and .Net, however, any language can be targeted by writing rules that identify patterns for that language or platform.

`csa` is entirely data driven using rules comprised of patterns that are first written in `yaml` and then compiled in the `csa` command-line executable. The rule system is flexible and can scan any type of written text, including source code, configuration files, and xml files. Basically, if the file is human-readable text, a rule can be devised that scans the file.

The matching of patterns and lines of application code require millions of pattern comparisons for each portfolio. To ensure performance of scans, `csa` is built to operate in a highly parallel manner. It is built in the Go language, which produces native-code executables for OSX, Windows, and Linux. `csa` will saturate all the CPUs of its host. Accordingly, `csa` benefits from running on multi-CPU machines, conversely, it suffers if it does not have multiple CPUs. **We recommend at least a 4-core (8 CPU) machine with 16 gig of RAM.**

The patterns are used to perform global scans of all application files, recursively in the directory specified on the command line. The rules are meant to be curated, and over time will change to adapt to the patterns found on cloud migration and containerization engagements. The intention is to create a single composite score that can be used assess cloud suitability, but also, to provide insights into an applications readiness for containerization.

Binaries/scripts

`csa` binaries run on the following platforms:

Executable	Platform
<code>csa.exe</code>	Windows
<code>csa-l</code>	Linux
<code>csa</code>	OSX

Installation

Download from here:

<https://github.com/vmware-samples/cloud-suitability-analyzer/releases>

There is no real installation process. It is just a matter of deciding on a home directory and copying the files in the `csa` distribution to that directory.

Setting up environment

To effectively use `csa` from the command-line, it will be helpful not to type in the full path every time. So include `csa`'s location in your path.

Adding the path on Linux

Change to your home directory.

```
cd $HOME
```

Open the `.bashrc` file with a text editor.

Add the following line to the file. Replace the with the location directory of `csa`

```
export PATH=<csa directory>:$PATH
```

Save the file and exit.

Use the source command to force Linux to reload the `.bashrc` file which normally is read only when you log in each time.

```
source .bashrc
```

Adding the path on OSX

Change to your home directory.

```
cd $HOME
```

Open the `.bash_profile` file with a text editor.

Add the following line to the file. Replace the with the location directory of `CSA`

```
export PATH=<csa directory>:$PATH
```

Save the file and exit.

Use the source command to force Linux to reload the `.bashrc` file which normally is read only when you log in each time.

```
source .bash_profile
```

Adding path on Windows

[Instructions to change your PATH on Windows 10](#)

File handles

If you are experiencing errors such as `Too many open files` or `Unable to open database file` you need to increase you open files `/maxfiles` ulimit.

If you are attempting to run `csa` on a large directory or set of directories this limit need to be set very high.

MAC OSX Sierra/High Sierra

LIMITED TEMPORARY FIX

Set ulimit to 20000

```
ulimit -n 20000
```

Note: This fix will only live as long as the current shell!

Persistent Fix (*more flexible...lets you set the limit higher. Requires Reboot!*)

1. You have to create a file in your root Library directory. Specifically =>

```
/Library/LaunchDaemons/limit.maxfiles.plist
```

Note: ensure (owner: root:wheel, mode: 0644) see steps below

2. Place the following into the file (set the limits as your desire). The settings below were tested against a portfolio of 36317 files with 7,331,920 lines of code.

```
<?xml version="1.0" encoding="UTF-8"?>
<!DOCTYPE plist PUBLIC "-//Apple//DTD PLIST 1.0//EN"
    "http://www.apple.com/DTDs/PropertyList-1.0.dtd">
<plist version="1.0">
  <dict>
    <key>Label</key>
    <string>limit.maxfiles</string>
    <key>ProgramArguments</key>
    <array>
      <string>launchctl</string>
      <string>limit</string>
      <string>maxfiles</string>
      <string>262144</string>
      <string>524288</string>
    </array>
    <key>RunAtLoad</key>
    <true/>
    <key>ServiceIPC</key>
    <false/>
  </dict>
</plist>
```

3. Make sure to set the permissions on the file correctly

```
=> sudo chown root:wheel /Library/LaunchDaemons/limit.maxfiles.plist
=> sudo chmod 0644 /Library/LaunchDaemons/limit.maxfiles.plist
```

4. Reboot your machine

Getting help from `csa`

`csa` has several major operating modes with thier own associated commands. You can see a list of the command by using `CSA help`

Command	Description
bins	Controls the creation of bins, which group similar applications together
rules	Add, edit, delete rules in the field, no need to rebuild executable.
naturalize	Future feature design to use machine learning to find bounded contexts
git	Run git forensics reports
search	Access the indexed search capabilities from the command line
analyze	Default command, scan a directory tree and apply rules
ui	Launch a local web server listening at localhost:3001

If you want help on any of these commands simple type `CSA help` and the command name, such as:

```
csa help rules
```

Cloning portfolios

`csa` expects to find a single application per sub-directory, if there are additional application in directory beneath the top directory, they will be considered as one application. This behavior can be controlled using configuration files. See below.

Using configuration files

Configuratoin files give you full control over how `csa` processes your application portfolio.

The table below describes the settings that are available:

Setting	Description
runName	Specify a run number, . lets the number be set by <code>CSA</code>

applications	A collection of application meta-data
Name	The name of the application, overrides directory name
Path	Directory where your application exists
business-domain	The domain or the department/region of the application
business-value	A number that indicates the value of the app to the business
dir-exclude-regex	A regex that describes directories that should be ignored
include-file-regex	A regex that includes files from processing
exclude-file-regex	A regex that excludes files from processing

Sample file

```
{
  "runName": ".",
  "applications": [
    {
      "Name": "App1",
      "Path": "/Users/swoods/pvtl/portfoliosmall-shortNames/App1",
      "business-domain": "",
      "business-value": 0,
      "dir-exclude-regex":
"^(\\.\\.\\.|target|classes|bin|test|node_modules|eclipse|out)$",
      "include-file-regex": ".*",
      "exclude-file-regex": "^(\\.\\.\\.|
(exe|png|tiff|tif|gif|jpg|jpeg|bmp|dmg|mpeg)|[\\.\\.\\.\\.|CSA-config\\.](yaml|yml|json))$"
    },
    {
      "Name": "App8",
      "Path": "/Users/swoods/pvtl/portfoliosmall-shortNames/App8",
      "business-domain": "",
      "business-value": 0,
      "dir-exclude-regex":
"^(\\.\\.\\.|target|classes|bin|test|node_modules|eclipse|out)$",
      "include-file-regex": ".*",
      "exclude-file-regex": "^(\\.\\.\\.|
(exe|png|tiff|tif|gif|jpg|jpeg|bmp|dmg|mpeg)|[\\.\\.\\.\\.|csa-config\\.](yaml|yml|json))$"
    }
  ]
}
```

Scoring system

Think of the scoring system as a measurement of relative effort to remediate an application to cloud-readiness. We use three loosely applied scales aligned with how often we expect to find a particular pattern in an applications source code.

Occurance	Score Range
Once per application	100-1000
Once per file	10 - 100
Multiple times per file	1 - 10

If the finding is really a positive, such as the discovery of spring boot pattern, then we make the number a negative. Since all scores are subtracted from a perfect score of 10, a negative score is essentially a positive.

For each application, once we add up the counts multiplied by the score we typically find a very wide range of scores between applications. Some may score a 50, while others may score 30,000. This stems from the fact that scoring is driven by lines of code per file. File size in software follows a `log-normal` distribution. So if we count anything related to file size, we will get a log-normal distribution. It looks like this:




Figure 1: Log-normal distribution

If we want to take the average, the median, or the standard deviation we need a normal distribution. A normal distribution looks like this:




Figure 2: Normal distribution

It's shaped like a bell and is sometimes called a bell-curve. Ever have a teacher that said they would grade you on a curve? This is that curve. Many human phenomena share this distribution. Take a room full of randomly selected people and ask them their height, you'll get a normal curve.

How do we reshape our curve? Since it is `log-normal`, really normal with a skew, we need a way to reshape it. This is very simple, we just take the base 10 logarithm of each number, in this case our total effort score.

###Default Model As discussed earlier, `csa`'s scoring system is externalized into a `yaml` files. The model also includes a set of thresholds to suggest depositions. These thresholds are experimental for the moment.

```
name: Default
#--- Although other models can be build, there always has to be
#    a Default model
max-score: 10
ranges:
  - type: sloc
    #--- Valid types include:
    #    sloc: Software Lines of Code
    #    raw: Raw score
    #    bv: Business value of app
    start: "0"
    end: int.max
    #--- describe a continuum of range bins
    ranges:
      #--- Bin description
      #    For any raw score between 0 and 100, Deploy to TAS, regardless
      #    of buiness score
```

```

- type: raw #--- raw score bin range
start: int.min #    start of raw bin
end: "100" #    end of raw bin
ranges:
  - type: bv #--- start of business value bin range
    start: float.min #    start of bv range
    end: float.max #    end of bv range
    outcome: #
      calculate: true
      #--- expression can be a complex formula that is based upon
      #    a combination of
      expression: max_score - log(10,raw_score)
      recommendation: Deploy to TAS

#--- Bin description
#    For any raw score between 101 and 100000
#    If BV less than 5, Rehost to TKG
#    If BV more than 5, Refactor to TAS
- type: raw
start: "101"
end: "10000"
ranges:
  - type: bv
    start: float.min
    end: "5.00"
    outcome:
      calculate: true
      expression: max_score - log(10,raw_score)
      recommendation: Rehost to TKG
  - type: bv
    start: "5.01"
    end: float.max
    outcome:
      calculate: true
      expression: max_score - log(10,raw_score)
      recommendation: Refactor to TAS

#--- Bin description
#    For any raw score between 10001 and 10000000
#    If BV less than 5, Rehost to TKG
#    If BV more than 5, Refactor to TAS
- type: raw
start: "10001"
end: "10000000"
ranges:
  - type: bv
    start: float.min
    end: "5.00"
    outcome:
      calculate: true
      expression: max_score - log(10,raw_score)

```



```

        recommendation: Rehost to TKG
- type: bv
  start: "5.01"
  end: flt.max
  outcome:
    calculate: true
    expression: max_score - log(10,raw_score)
    recommendation: Refactor to TAS

#--- Bin description
#   For any raw score greater than 10000001
#   If BV less than 5, Rehost to TKG
#   If BV more than 5, Refactor to TAS

- type: raw
  start: "10000001"
  end: int.max
  ranges:
    - type: bv
      start: flt.min
      end: "5.00"
      outcome:
        calculate: true
        expression: max_score - log(10,raw_score)
        recommendation: Rehost to TKG
    - type: bv
      start: "5.01"
      end: flt.max
      outcome:
        calculate: true
        expression: max_score - log(10,raw_score)
        recommendation: Refactor to TAS

```

Adding rules

An important design requirement for `csa` was the ability to change rules in the field, without the need to recompile the executable. This requirement is driven by the realization that many customer may have in-house libraries that have `wrapper` classes and functions to simplify the use of other frameworks. As such, these wrapper classes may hide critical patterns. With this capability, those internal libraries can be scanned first and then the rules may be augmented to look for additional patterns. The following process details the steps required to do this.

1. Export the rules currently contained inside the `CSA` executable

```
csa rules export --output-dir=./rules
```

2. Add whatever rules you want to add to the rules directory that was just created. You may edit existing rules, as well.

3. Run the following command

```
csa rules import --rules-dir=./rules
```

4. Open a shell window (bash, MingW, git bash, powershell, etc...)

5. Run `csa help` with no parameters and you'll see usage instructions.

6. To target the current directory and analyze it simply type

```
csa -p .
```

The `-p` tell `csa` to treat each sub directory as a stand-alone application. Otherwise, the sub directories score will be rolled into a single application. In other words, the sub-directories are considered parts of a single application.

7. To target a directory with source code simple run

```
csa -p <path> or csa analyze -p <path>
```

In most usages it is expected that the user has `git` cloned multiple applications into a single directory and therefore each of those sub-directories is a single application.

Targeting an ear/war/jar

If no source code is available, you can decompile the ear, war, or jar files. To do so, you'll need the jar file `fernflower.jar` that is bundled in the `csa` download. We suggest putting the jar in the same directory as the `csa` executable, but this can be overridden with the `--fern-jar-path` flag.

1. Run `csa` and provide the fully qualified path to the "jar"

```
csa analyze -p ~/resteasy-spring-2.3.8.Final-redhat-3.jar
```

Tool output

`csa` provides various useful outputs as it processes applications. These can be useful in understanding the results of the scan.

- Command `csa` is executing
- Critical directories
- Number of applications discovered
- Total files found in each application sub-directory
- Percent progress for each scan
- Software Lines of Code (SLOC) summary

While the scan occurs, `csa` is also loading a Sqlite database called `csa.db`. The default location is in the same directory in which the `csa` executable is located. When you first download `csa`, the `csa.db` does not exist, it is created the first time you run `csa`, your first run will be labeled `run 1` and each subsequent run will be incremented. You can have as many run's as you like in `csa.db`. If you want to start with an empty `csa.db` you can delete or rename the file.

NOTE: If you download a new version of `csa` you will need to delete/rename the current `csa.db` to have any new rules appear in `csa`.

Rules

What is a Rule? A rule is in simplest terms a description of something that you want `csa` to detect. This description is structured so that `csa` can easily understand it but is designed to be flexible and extensible.

Important Note : Rules and analysis data are intentionally ephemeral from the perspective of `csa` and the baseline code. *If you update rules, as discussed above, you'll want to keep them in some version controlled system, such as git. They represent valuable insight into your portfolio's profile.*

Understanding rules

Rule model

Attribute	Type	Description	Required (y/n)	Default
Name	string	The name of the rule. Can be meaningful or not but must be unique! And must match the name of the yaml file.	Y	
FileType	string	The file extension the rule will target. I.E. <code>java</code> for <code>.java</code> files! Value should not include the dot (period). This can also be a regular expression. I.E. <code>xml[li]</code> would match both <code>xml</code> and <code>xmi</code> files	N	Rule apply all files no value is specified
Target	enum	This is the target of the rule. Valid values: <code>File</code> , <code>Line</code> . <code>File</code> = rule will apply to filenames only. <code>Line</code> = rule will be applied against every line of content within the file.	Y	
Type	enum	This specifies the type or behavior of the rule. Valid values: <code>regex</code> , <code>simple-text</code> , <code>simple-text-ci</code> , <code>starts-with</code> , <code>starts-with-ci</code> , <code>ends-with</code> , <code>ends-with-ci</code> , <code>contains</code> , <code>contains-ci</code>	Y	
DefaultPattern	string	Pattern with a	N	

		placeholder (%s) for substitution of "Pattern" values. I.E. "[.]%s[]". This does not only apply to Regex rules but can also be used for others like a StartsWith such as 'org.json.%s'		
Advice	string	Any advice on how to remediate this finding for cloud compatibility. This value is used if the specific pattern does not have advice.	N	
Score	int	A value indicating how this finding impacts cloud compatibility. At this time we have not settled on a scoring model so ...	N	
Category	string	The category of the rule. Simply a text marker to allow for grouping during analysis in csa. I.E. For the API rules this contains the API name	N	
Criticality	enum	A t-shirt size of the impact of the finding. Valid values: High, Medium, Low. Used for dashboard in csa	N	
Tags	array of Tag objects	Tags is a collection (0-n) of string values that can be used for grouping/slicing/etc... during analysis in csa	N	
Recipes	array of	Recipes is a collection (0-n) of URI values pointing at	N	

	Recipe objects	applicable recipes to aid in remediation of the finding		
Patterns	array of Pattern objects	Patterns contains the patterns (1-n) that will be used to match against filenames/line data and result in findings	Y (at least 1)	

Pattern model

Attribute	Type	Description	Required (y/n)	Default
Value	string	This is the actual pattern value! It will be substituted into or the default pattern or the overriding pattern.	Y	
Type	enum	This specifies the type or behavior of the pattern. Overrides the rule type. Valid values: regex, simple-text, simple-text-ci, starts-with, starts-with-ci, ends-with, ends-with-ci, contains, contains-ci	Y	
Pattern	string	Pattern with a placeholder (%s) for substitution of Value. I.E. "[.]%s[()]" This does not only apply to Regex rules but can also be used for others like a StartsWith such as 'org.json.%s'	N	
Advice	string	Any advice on how to remediate this finding for cloud compatibility.	N	

		Overrides any advice provided at the rule level.		
Score	int	A value indicating how this finding impacts cloud compatibility. At this time we have not settled on a scoring model so ...Overrides any score provided at the rule level.	N	
Criticality	enum	A t-shirt size of the impact of the finding. Valid values: High, Medium, Low. Used for dashboard in csa. Overrides any Criticality provided at the rule level.	N	
Tags	array of Tag objects	Tags is a collection (0-n) of string values that can be used for grouping/slicing/etc... during analysis in csa. Overrides any tags provided at the rule level.	N	

Tag model

Attribute	Type	Description	Required (y/n)	Default
Value	string	the string you are tagging the rule or pattern with	N	

Recipe model

Attribute	Type	Description	Required (y/n)	Default
URI	string	A <code>uri</code> for the recipe to resolve the finding	N	

Example Rules (yaml)

Line level Regex

This is the default annotations rule. It is only be applied against `.java` files, will be matched against every line in the file and detects this use of the annotations listed under patterns.

```
name: annotations
filetype: java
target: line
type: regex
defaultpattern: ^.*@%s$
criticality: medium
tags:
  - value: annotations
patterns:
  - value: DeclareRoles
  - value: DenyAll
  - value: PermitAll
  - value: RolesAllowed
  - value: RunAs
  - value: Stateless
  - value: Stateful
  - value: MessageDriven
  - value: Entity
  - value: Init
  - value: Remove
  - value: ActivationConfigProperty
  - value: Local
  - value: Remote
  - value: LocalHome
  - value: RemoteHome
  - value: TransactionManagement
  - value: TransactionAttribute
  - value: PostActivate
  - value: PreTASsivate
```

File Target

This rule only gets applied against java files and detects the presence of pattern named under patterns section.

```
name: java-iop
filetype: java$
target: line
type: regex
advice: Move to cloud friendly alternatives
defaultpattern: "^.*[.]%s[ (.)\\.]"
effort: 100
readiness: 6
category: iop
tags:
  - value: api
  - value: protocol
  - value: ejb
```

```
- value: non-standard
patterns:
- value: PortableRemoteObject
- value: CodecFactory
- value: CodecOperations
- value: TransactionService
- value: ServiceContext
- value: TaggedComponent
- value: TaggedProfile
```

Rules management

Exporting

So, now you understand rules. What rules come by default? Or what do the current set of rules that `csa` is using look like? Let's export them!

Run the `csa` rules command with the export sub-command. By default rules will export to the default output-dir. You can override where they will go with the `--output-dir` flag or the `--rules-dir` flag. By default they will be exported as yaml with each rule in a separate file. If you prefer json or only working with a single file there are command flags to control this behavior. run ```csa help rules export``` for details.

```
usage: csa rules export [<flags>] [<name>]
```

```
export rule(s) from the database
```

```
==> csa rules export
DBEngine: sqlite Name: csa.db Version: 3.23.1
Successfully exported [47] rules @ [csa-reports/rules]
```

```
==> csa rules export
DBEngine: sqlite Name: csa.db Version: 3.23.1
Successfully exported [47] rules @ [csa-reports/rules]
```

Creating/Updating/Importing rules

So, you can now see the rules that come by default. You want to change one, edit the file and update the rule. You want to create one, create a new file with the appropriate structure or add the rule to an existing file. Then run the `csa rules import` command. By default all rules in the `--rules-dir` will be imported or you can specify a rule name as an argument to the command. There are command flags to control directory where rules will be read and whether rules will be replaced or updated. Run `csa help rules import` for details.

```
usage: csa rules import [<flags>] [<name>]
```

```
import rule(s) into the database. By default rules will be added/updated rather than
replace existing
```

```
==> csa rules import --rules-dir=csa-reports/rules
DBEngine: sqlite Name: csa.db Version: 3.23.1
Successfully imported [47] rule(s) found @[csa`csa`-reports/rules]
```


Note: If importing more than one rule for file ==> If file format is yaml follow the standard yaml multi-document format of separating documents with `---`. If file format is json then just put the rule (object) in the file as a distinct object. Json really doesn't support more than one top level object in a file but that's ok! :). For example of how to create a multi-doc file run the export with the flag to create a single file and review!

Deleting/Removing

You have a rule you don't want anymore. Or, for some reason, you want a clean slate...

Delete a rule

```
usage: csa rules delete <name>
```

delete a rule in the database

```
==> csa rules delete annotations
DBEngine: sqlite Name: csa.db Version: 3.23.1
Deleting rule [annotations]...done!
```

Note: If the rule is found you will receive an indication it is deleted. If it is not found...you won't see any indication other than a clean exit(0) from the CLI

Delete All Rules (caution advised!)

```
usage: csa rules delete-all
```

delete all rules in the database!

```
==> csa rules delete-all
DBEngine: sqlite Name: csa.db Version: 3.23.1
Delete All Rules! Are you sure(y/n)? y
All Rules Successfully Deleted!
```

Note: Rule 'filenames' are unimportant and have no bearing on rule behavior and are only important to the OS to disambiguate one file from another. Rule 'names' are only important from the perspective of they must be unique.

Application Archetypes

Bucketing of applications by tags

All rules in `csa` have any number of `tags` associated with them. A `tag` just associates a concept with the rule, such as `jni` or `corba`. At a higher level, we can think of a group of tags can further identify an architecture archetype or a bucket of similar applications. `csa` uses a single `yaml` file (`bins/bins.yaml`) to describe the archetypes as seen in the excerpt below:

```
name: TKG
tags:
- name: Docker
  type: OR
- name: stateful
  type: AND
- name: javaee
```

```
type: AND
- name: fullprofile
  type: AND
- name: jni
  type: OR
- name: nonstandard-protocol
  type: OR
- name: corba
  type: OR
```

```
name: TAS
tags:
- name: webprofile
  type: OR
- name: spring
  type: OR
- name: spring-boot
  type: OR
- name: webcontainer
  type: OR
- name: rest
  type: OR
- name: jar
  type: OR
```

```
.
.
.
```